

## Lab Assignment – 2

### Implement Ridge Regularization and the Lasso Regularization in Python

**Lasso Regularization:** In lasso regularization, a penalty term is added to the standard loss function of the model, which is typically the sum of squared errors (for regression) or the log-likelihood (for classification). This penalty term is the sum of the absolute values of the coefficients of the model multiplied by a regularization parameter, often denoted as  $\lambda$  (lambda). The formula for the loss function with lasso regularization can be expressed as:

Loss function +  $\lambda$  \* (sum of absolute values of coefficients)

**Ridge Regularization:** In ridge regularization, a penalty term is added to the standard loss function of the model, typically the sum of squared errors (for regression) or the log-likelihood (for classification). This penalty term is the sum of the squared magnitudes of the coefficients of the model multiplied by a regularization parameter, often denoted as  $\lambda$  (lambda). The formula for the loss function with ridge regularization can be expressed as:

Loss function +  $\lambda$  \* (sum of squared magnitudes of coefficients)

Code:-

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import scale
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge, RidgeCV, Lasso, LassoCV
from sklearn.metrics import mean_squared_error
```

```
df = pd.read_csv('Hitters.csv').dropna()
df.drop(df.columns[[0]], axis = 1, inplace = True)
df.info()
dummies = pd.get_dummies(df[['League', 'Division', 'NewLeague']])
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 263 entries, 1 to 321
Data columns (total 20 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   AtBat       263 non-null    int64
 1   Hits        263 non-null    int64
 2   HmRun       263 non-null    int64
 3   Runs        263 non-null    int64
 4   RBI         263 non-null    int64
 5   Walks       263 non-null    int64
 6   Years       263 non-null    int64
 7   CAtBat      263 non-null    int64
 8   CHits       263 non-null    int64
 9   CHmRun      263 non-null    int64
10   CRuns       263 non-null    int64
11   CRBI        263 non-null    int64
12   CWalks      263 non-null    int64
13   League      263 non-null    object
14   Division    263 non-null    object
15   PutOuts     263 non-null    int64
16   Assists     263 non-null    int64
17   Errors      263 non-null    int64
18   Salary      263 non-null    float64
19   NewLeague   263 non-null    object
dtypes: float64(1), int64(16), object(3)
memory usage: 43.1+ KB

```

```

y = df.Salary
# Drop the column with the independent variable (Salary), and columns for which we created dummy variables
X_ = df.drop(['Salary', 'League', 'Division', 'NewLeague'], axis = 1).astype('float64')
# Define the feature set X.
X = pd.concat([X_, dummies[['League_N', 'Division_W', 'NewLeague_N']]], axis = 1)
X.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 263 entries, 1 to 321
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   AtBat           263 non-null   float64
1   Hits            263 non-null   float64
2   HmRun           263 non-null   float64
3   Runs            263 non-null   float64
4   RBI             263 non-null   float64
5   Walks           263 non-null   float64
6   Years           263 non-null   float64
7   CAtBat          263 non-null   float64
8   CHits           263 non-null   float64
9   CHmRun          263 non-null   float64
10  CRuns           263 non-null   float64
11  CRBI            263 non-null   float64
12  CWalks          263 non-null   float64
13  PutOuts         263 non-null   float64
14  Assists         263 non-null   float64
15  Errors          263 non-null   float64
16  League_N        263 non-null   uint8
17  Division_W      263 non-null   uint8
18  NewLeague_N     263 non-null   uint8
dtypes: float64(16), uint8(3)
memory usage: 35.7 KB

```

```
alphas = 10**np.linspace(10,-2,100)*0.5
alphas
```

```
array([5.00000000e+09, 3.78231664e+09, 2.86118383e+09, 2.16438064e+09,
       1.63727458e+09, 1.23853818e+09, 9.36908711e+08, 7.08737081e+08,
       5.36133611e+08, 4.05565415e+08, 3.06795364e+08, 2.32079442e+08,
       1.75559587e+08, 1.32804389e+08, 1.00461650e+08, 7.59955541e+07,
       5.74878498e+07, 4.34874501e+07, 3.28966612e+07, 2.48851178e+07,
       1.88246790e+07, 1.42401793e+07, 1.07721735e+07, 8.14875417e+06,
       6.16423370e+06, 4.66301673e+06, 3.52740116e+06, 2.66834962e+06,
       2.01850863e+06, 1.52692775e+06, 1.15506485e+06, 8.73764200e+05,
       6.60970574e+05, 5.00000000e+05, 3.78231664e+05, 2.86118383e+05,
       2.16438064e+05, 1.63727458e+05, 1.23853818e+05, 9.36908711e+04,
       7.08737081e+04, 5.36133611e+04, 4.05565415e+04, 3.06795364e+04,
       2.32079442e+04, 1.75559587e+04, 1.32804389e+04, 1.00461650e+04,
       7.59955541e+03, 5.74878498e+03, 4.34874501e+03, 3.28966612e+03,
       2.48851178e+03, 1.88246790e+03, 1.42401793e+03, 1.07721735e+03,
       8.14875417e+02, 6.16423370e+02, 4.66301673e+02, 3.52740116e+02,
       2.66834962e+02, 2.01850863e+02, 1.52692775e+02, 1.15506485e+02,
       8.73764200e+01, 6.60970574e+01, 5.00000000e+01, 3.78231664e+01,
       2.86118383e+01, 2.16438064e+01, 1.63727458e+01, 1.23853818e+01,
       9.36908711e+00, 7.08737081e+00, 5.36133611e+00, 4.05565415e+00,
       3.06795364e+00, 2.32079442e+00, 1.75559587e+00, 1.32804389e+00,
       1.00461650e+00, 7.59955541e-01, 5.74878498e-01, 4.34874501e-01,
       3.28966612e-01, 2.48851178e-01, 1.88246790e-01, 1.42401793e-01,
       1.07721735e-01, 8.14875417e-02, 6.16423370e-02, 4.66301673e-02,
       3.52740116e-02, 2.66834962e-02, 2.01850863e-02, 1.52692775e-02,
       1.15506485e-02, 8.73764200e-03, 6.60970574e-03, 5.00000000e-03])
```

```
ridge = Ridge()
coefs = []
for a in alphas:
    ridge.set_params(alpha = a)
    ridge.fit(X, y)
    coefs.append(ridge.coef_)

np.shape(coefs)
```


```
(100, 19)
```

```

ax = plt.gca()
ax.plot(alphas, coefs)
ax.set_xscale('log')
plt.axis('tight')
plt.xlabel('alpha')
plt.ylabel('weights')

```

```
Text(0, 0.5, 'weights')
```



```

# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.5, random_state=1)

```

```

ridge2 = Ridge(alpha = 4)
ridge2.fit(X_train, y_train) # Fit a ridge regression on the training data
pred2 = ridge2.predict(X_test) # Use this model to predict the test data
print(pd.Series(ridge2.coef_, index = X.columns)) # Print coefficients
print(mean_squared_error(y_test, pred2)) # Calculate the test MSE

```

```

AtBat          -1.847687
Hits           4.359634
HmRun          -5.277385
Runs           -0.166042
RBI            4.247790
Walks          3.492435
Years          9.174059
CAtBat         -0.595198
CHits          2.140742
CHmRun         2.913484
CRuns          0.271946
CRBI           -0.591043
CWalks         0.171610
PutOuts        0.418292
Assists        0.445158
Errors         -5.507812
League_N       67.441132
Division_W     -100.612196
NewLeague_N    -21.153561
dtype: float64
115325.04539923287

```

```
ridge3 = Ridge(alpha = 10**10)
ridge3.fit(X_train, y_train) # Fit a ridge regression on the training data
pred3 = ridge3.predict(X_test) # Use this model to predict the test data
print(pd.Series(ridge3.coef_, index = X.columns)) # Print coefficients
print(mean_squared_error(y_test, pred3)) # Calculate the test MSE
```

```
AtBat      3.610286e-04
Hits       1.285729e-04
HmRun      1.687679e-05
Runs       6.720858e-05
RBI        8.010120e-05
Walks      6.606546e-05
Years      1.078452e-05
CAtBat     6.880730e-03
CHits      2.095391e-03
CHmRun     2.864823e-04
CRuns      1.071697e-03
CRBI       1.189446e-03
CWalks     7.448325e-04
PutOuts    8.437728e-04
Assists    -4.683932e-06
Errors     1.266266e-06
League_N   -7.484549e-08
Division_W -5.136370e-07
NewLeague_N -5.969551e-08
dtype: float64
165016.14024085485
```

```

ridge2 = Ridge(alpha = 0)
ridge2.fit(X_train, y_train) # Fit a ridge regression on the training data
pred = ridge2.predict(X_test) # Use this model to predict the test data
print(pd.Series(ridge2.coef_, index = X.columns)) # Print coefficients
print(mean_squared_error(y_test, pred)) # Calculate the test MSE

```

```

AtBat      -1.821115
Hits       4.259156
HmRun     -4.773401
Runs      -0.038760
RBI        3.984578
Walks      3.470126
Years      9.498236
CAtBat    -0.605129
CHits      2.174979
CHmRun     2.979306
CRuns      0.266356
CRBI       -0.598456
CWalks     0.171383
PutOuts    0.421063
Assists    0.464379
Errors     -6.024576
League_N   133.743163
Division_W -113.743875
NewLeague_N -81.927763
dtype: float64
116690.46856662024

```

```

ridgecv = RidgeCV(alphas = alphas, scoring = 'neg_mean_squared_error')
ridgecv.fit(X_train, y_train)
ridgecv.alpha_

```

```

378231.66377731453

```

```

ridge4 = Ridge(alpha = ridgecv.alpha_)
ridge4.fit(X_train, y_train)
mean_squared_error(y_test, ridge4.predict(X_test))

```

```

113938.33161417228

```




```
ridge4.fit(X, y)
pd.Series(ridge4.coef_, index = X.columns)
```

```
AtBat      0.212638
Hits       0.582154
HmRun      0.021942
Runs       0.333584
RBI        0.191530
Walks      0.507784
Years     -0.007532
CAtBat     -0.311717
CHits      0.800669
CHmRun     0.165121
CRuns      0.715549
CRBI       0.567190
CWalks    -0.025906
PutOuts    0.290130
Assists    0.154169
Errors    -0.067756
League_N   0.006936
Division_W -0.022272
NewLeague_N 0.004730
dtype: float64
```

```
lasso = Lasso(max_iter = 10000)
coefs = []
for a in alphas:
    lasso.set_params(alpha=a)
    lasso.fit(scale(X_train), y_train)
    coefs.append(lasso.coef_)

ax = plt.gca()
ax.plot(alphas*2, coefs)
ax.set_xscale('log')
plt.axis('tight')
plt.xlabel('alpha')
plt.ylabel('weights')
```

Text(0, 0.5, 'weights')



```
lassocv = LassoCV(alphas = None, cv = 10, max_iter = 100000)
lassocv.fit(X_train, y_train)
lasso.set_params(alpha=lassocv.alpha_)
lasso.fit(X_train, y_train)
mean_squared_error(y_test, lasso.predict(X_test))
```

118255.15053121911

```
pd.Series(lasso.coef_, index=X.columns)
```

AtBat	0.186425
Hits	0.000000
HmRun	0.000000
Runs	0.000000
RBI	0.000000
Walks	0.000000
Years	-0.000000
CAtBat	-0.508102
CHits	1.424662
CHmRun	0.000000
CRuns	0.597670
CRBI	0.548273
CWalks	0.404799
PutOuts	0.456655
Assists	0.048838
Errors	-0.000000
League_N	0.000000
Division_W	-0.000000
NewLeague_N	0.000000

dtype: float64